

# Logboek

Kasper Notebomer

9/14/2021

Setup of the libraries.

```
# Libraries
library("ggplot2")
```

```
## Registered S3 methods overwritten by 'tibble':
##   method      from
##   format.tbl  pillar
##   print.tbl   pillar
```

```
## Warning: replacing previous import 'vctrs::data_frame' by 'tibble::data_frame'
## when loading 'dplyr'
```

```
library("kableExtra")
library("factoextra")
```

```
## Welcome! Want to learn more? See two factoextra-related books at https://goo.gl/ve3WBa
```

```
library("cowplot")
library("gridExtra")
library("dplyr")
```

```
##
## Attaching package: 'dplyr'
```

```
## The following object is masked from 'package:gridExtra':
```

```
##
##   combine
```

```
## The following object is masked from 'package:kableExtra':
```

```
##
##   group_rows
```

```
## The following objects are masked from 'package:stats':
```

```
##
##   filter, lag
```

```
## The following objects are masked from 'package:base':
```

```
##
##   intersect, setdiff, setequal, union
```

```
library("scales")
library(ggpubr)
```

```
##
## Attaching package: 'ggpubr'

## The following object is masked from 'package:cowplot':
##
##   get_legend
```

```
library(ggfortify)
library(tidyr)
library(purrr)
```

```
##
## Attaching package: 'purrr'

## The following object is masked from 'package:scales':
##
##   discard
```

```
library(ggcorrplot)
library(reshape2)
```

```
##
## Attaching package: 'reshape2'

## The following object is masked from 'package:tidyr':
##
##   smiths
```

```
library(e1071)
library(Hmisc)
```

```
## Loading required package: lattice

## Loading required package: survival

## Loading required package: Formula
```

```
##
## Attaching package: 'Hmisc'

## The following object is masked from 'package:e1071':
##
##   impute
```

```
## The following objects are masked from 'package:dplyr':
##
##   src, summarize
```

```
## The following objects are masked from 'package:base':
##
##   format.pval, units
```

# Introduction

The data was provided by Tsjerk Wassenaar of the RuG. The data set contains data on membrane composition and characteristics. The data set is not publicly available. The data set also does not have a publication linked to it.

Seeing as there is no paper linked to the data set I can only guess using the information that I did get that the data was gathered by making a membrane using specific variables and then measuring the resulting membranes to get variables like the thickness and compressibility. The data set contains 14 variables and 2843 different measurements. The goal is to use measurements to predict the composition of the membrane. Typically you would use independent variables to predict one dependent variable. In this data set this doesn't seem to be the case, seeing as the variables that are given as parameters are seen as class variables, this means that these would be considered the labels. The parameters are: Temperature, Sterol type, Sterol concentration, Other (phospho)lipids in membrane, Aliphatic tails, Saturation index, Phosphatidyl choline concentration and Ethanol concentration. So the biggest question that the EDA should answer is which of these variables is most interesting to use as the label, or could we even predict multiple of them using the given variables.

## Data exploration

## Data reading & codebook

To start off the data is read in using `read.csv()`. The data frame that is created by this step is turned into a tibble and then used to create a codebook. The codebook is a csv file that contains the abbreviation, type and the description of every variable in the data set.

```
# Read the data and transform it to a tibble
data <- read.csv("data/dataFrame_all_sims.csv",
                 sep = ",", header=TRUE,
                 stringsAsFactors=FALSE)
data <- as_tibble(data)
```

Here, the codebook is made of all of the abbreviations in the data frame, their meaning and their type.

```
# Create the codebook
codebook <- data.frame("Abbreviation" = colnames(data), "Type" = sapply(data, typeof),
  "Class" = sapply(data, class),
  "Description" = c("Temperature (Kelvin)",
    "Sterol type",
    "Sterol concentration (%)",
    "Other (phospho)lipids in membrane (headgroup)",
    "Aliphatic tails",
    "Saturation index (double bonds per tail)",
    paste("Phosphatidyl choline concentration",
      "(% of non-sterol lipids)"),
    "Ethanol concentration (% of solvent)",
    "Area per lipid (nm^2)",
    "Thickness (nm)", "Bending rigidity (kB T)",
    "Tilt angle (degrees)",
    "Z-order", "Compressibility (cN / m)",
  "Units" = c("Kelvin",
    ""
```

```

      "%",
      "",
      "",
      "",
      "",
      "",
      "",
      "nm^2",
      "nm",
      "kB T",
      "degrees",
      "",
      "cN / m"))

# Turn dataframe into a tibble
codebook <- as_tibble(codebook)
kable(codebook, caption = "Codebook") %>%
  kable_styling(latex_options = c("scale_down", "hold_position"))

```

Table 1: Codebook

Abbreviation	Type	Class	Description	Units
temperature	integer	integer	Temperature (Kelvin)	Kelvin
sterol.type	character	character	Sterol type	
sterol.conc	integer	integer	Sterol concentration (%)	%
other.phosph	character	character	Other (phospho)lipids in membrane (headgroup)	
tails	character	character	Aliphatic tails	
satur.index	double	numeric	Saturation index (double bonds per tail)	
PC.conc	integer	integer	Phosphatidyl choline concentration (% of non-sterol lipids)	
ethanol.conc	integer	integer	Ethanol concentration (% of solvent)	
APL	double	numeric	Area per lipid (nm <sup>2</sup> )	nm <sup>2</sup>
thickness	double	numeric	Thickness (nm)	nm
bending	double	numeric	Bending rigidity (kB T)	kB T
tilt	double	numeric	Tilt angle (degrees)	degrees
zorder	double	numeric	Z-order	
compress	double	numeric	Compressibility (cN / m)	cN / m

```

# Write codebook to .csv file
write.csv(codebook, "data/codebook.csv", row.names = FALSE)

```

The codebook, shown in table 1, seems to have all of the correct data and datatypes. To easily access the description based on the abbreviation the following function is used.

```

# Get the description by the abbreviation from the codebook
get_des_by_ab <- function(df, abbreviation){
  description <- df[df$Abbreviation == abbreviation,]$Description
  return(description)
}

```

## Structure & summary

Here I'll take a look at the structure and the five number summary of the data to look for any problems or discrepancies.

```
# Print the structure of the data
data.frame(variable = names(data),
           classe = sapply(data, class),
           first_values = sapply(data, function(x) paste0(head(x), collapse = ", ")),
           row.names = NULL) %>%
  kable(caption="Data structure") %>%
  kable_styling(latex_options = c("scale_down", "hold_position"))
```

Table 2: Data structure

variable	classe	first_values
temperature	integer	298, 298, 298, 298, 298, 298
sterol.type	character	chole, chole, chole, chole, chole, chole
sterol.conc	integer	20, 20, 20, 20, 20, 20
other.phosph	character	PE, PE, PE, PE, PE, PE
tails	character	PI, PI, PI, PI, PI, PI
satur.index	numeric	1, 1, 1, 1, 1, 1
PC.conc	integer	33, 50, 25, 100, 0, 75
ethanol.conc	integer	20, 20, 20, 20, 20, 20
APL	numeric	0.736911, 0.746226, 0.731945, 0.76963, 0.719509, 0.760091
thickness	numeric	3.62897268631964, 3.59699140829278, 3.6526367497855, 3.49605162681907, 3.71001820243157, 3.53890823219068
bending	numeric	10.0802, 10.1403, 10.2547, 10.0232, 10.4654, 9.86163
tilt	numeric	14.7578, 14.3075, 14.9777, 13.2029, 15.753, 13.7148
zorder	numeric	0.175499, 0.172084, 0.178494, 0.163349, 0.182191, 0.166807
compress	numeric	24.1821842679417, 24.3249489613324, 23.766069963793, 22.9607685634682, 23.1050322963544, 23.8604998172132

Looking at the structure of the data, shown in table 2, it seems like all of the columns have been read and have the right datatype.

```
# Get the five number summary of the selected columns
sum <- summary(data[c("temperature", "sterol.conc", "satur.index",
                     "PC.conc", "ethanol.conc", "APL", "thickness", "bending",
                     "tilt", "zorder", "compress")])

sum <- sub(".*:", "", sum)
sum[is.na(sum)] <- 0
rownames(sum) <- c("Minimum", "Q1", "Median", "Mean", "Q3", "Maximum", "Number of NA's")

# Print five number summary using kable
kable(sum, caption="Five number summary") %>%
  kable_styling(latex_options = c("scale_down", "hold_position"))
```

Table 3: Five number summary

	temperature	sterol.conc	satur.index	PC.conc	ethanol.conc	APL	thickness	bending	tilt	zorder	compress
Minimum	298.0	0.00	0.0000	0.00	0	0.4514	3.003	0.2025	6.774	0.02736	1.286
Q1	298.0	10.00	0.0000	25.00	5	0.6350	3.627	9.8579	14.490	0.18604	23.097
Median	298.0	20.00	0.5000	50.00	15	0.6885	3.812	12.1000	18.157	0.25860	31.346
Mean	305.2	16.55	0.6208	50.02	15	0.6807	3.934	22.0256	35.732	0.33269	48.714
Q3	298.0	30.00	1.0000	75.00	25	0.7394	4.238	25.5712	35.933	0.45559	43.348
Maximum	328.0	30.00	2.0000	100.00	30	0.8924	5.101	125.3820	227.813	0.93013	538.446
Number of NA's	0	0	0	0	0	104	103	104	104	153	103

Looking at the structure there doesn't seem to be immediate problem with the data, like the data being read as the wrong datatype. When looking at the five number summary, depicted in table 3, there do seem to be some rows that are missing values like the APL, thickness, bending and tilt values. These rows will be omitted for plotting and other research seeing as they can't be used. The maximum of the bending rigidity

and tilt angle seem to be too high, this probably means that these are outliers seeing as the rest of the values in the summary seem to be believable. To check this we'll make a histogram and a boxplot of these variables.

```
# Remove NA's
row.has.na <- apply(data, 1, function(x){any(is.na(x))})
nona_data <- na.omit(data)
row.has.na.2 <- apply(nona_data, 1, function(x){any(is.na(x))})
paste("Number of NA's before: ", sum(row.has.na),
      ", Number of NA's after:", sum(row.has.na.2), sep="")
```


```
## [1] "Number of NA's before: 153, Number of NA's after:0"
```

153 lines with NA's were omitted from the data set.

Here is a boxplot of all of the variables to check for any outliers.

```
nona_data %>%
  select("temperature", "sterol.conc", "satur.index",
         "PC.conc", "ethanol.conc", "APL", "thickness", "bending",
         "tilt", "zorder", "compress") %>%
  pivot_longer(., cols = c("temperature", "sterol.conc", "satur.index",
                           "PC.conc", "ethanol.conc", "APL", "thickness", "bending",
                           "tilt", "zorder", "compress"), names_to = "Variable", values_to = "Value") %>%
  ggplot(aes(x = Variable, y = Value)) +
  geom_boxplot() + theme(axis.text.x = element_text(angle = 90, vjust = 0.5, hjust=1)) +
  ggtitle("Figure 0: Histogram of the variables")
```

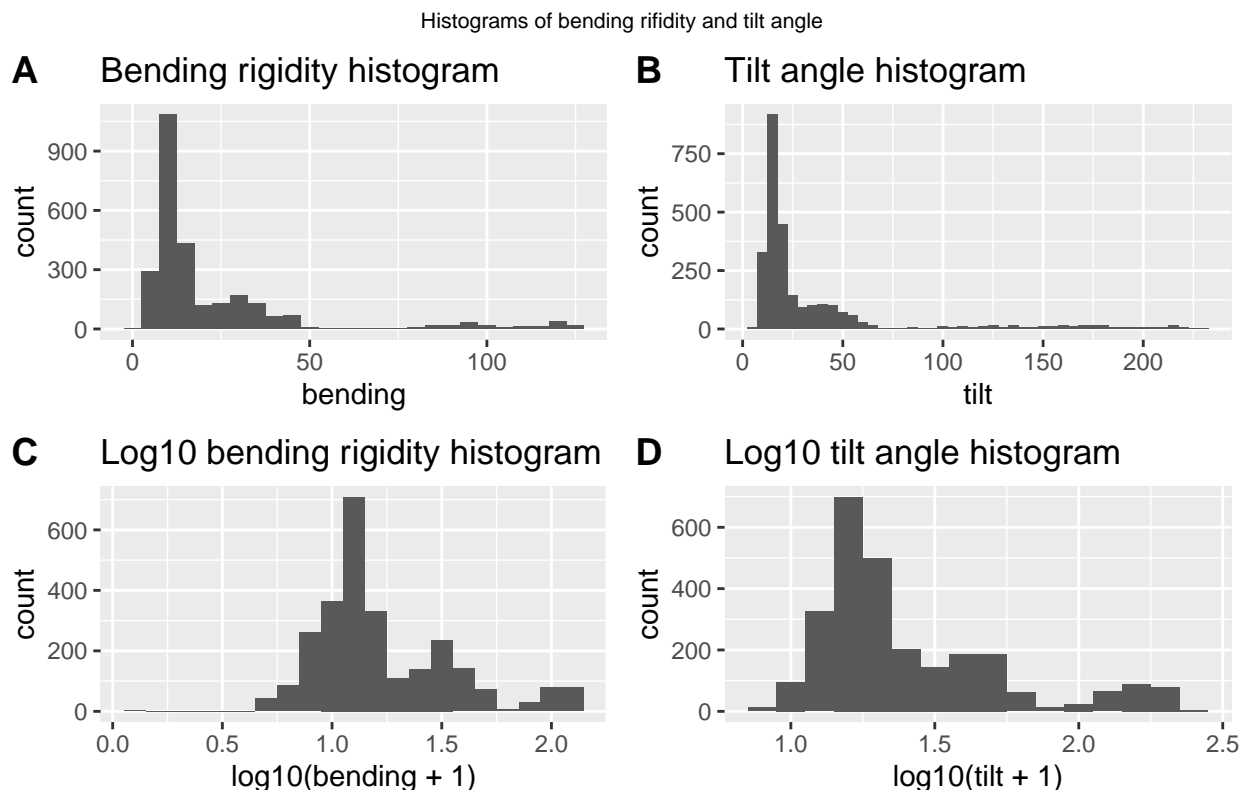
Figure 0: Histogram of the variables



The boxplot clearly shows that there are indeed a lot of outliers to be found in the dataset, especially in the bending rigidity, compressibility and tilt angle columns. Outliers are shown as black dots in the plot. Some variables also have very clean distributions, this is mostly because some variables are numerical but only contain a few different levels because of the experiment that they were gathered from.

```
# Create 4 histograms
p <- ggplot(nona_data, aes(x= bending)) +
  geom_histogram(binwidth=5) +
  ggtitle("Bending rigidity histogram")
p2 <- ggplot(nona_data, aes(x=tilt)) +
  geom_histogram(binwidth=5) +
  ggtitle("Tilt angle histogram")
p3 <- ggplot(nona_data, aes(x= log10(bending + 1))) +
  geom_histogram(binwidth=0.1) +
  ggtitle("Log10 bending rigidity histogram")
p4 <- ggplot(nona_data, aes(x=log10(tilt + 1))) +
  geom_histogram(binwidth=0.1) +
  ggtitle("Log10 tilt angle histogram")
plots <- ggarrange(p, p2, p3, p4,
  labels = c("A", "B", "C", "D"),
  ncol = 2, nrow = 2)
title <- expression(atop(bold("Figure 1:"),
  scriptstyle("Histograms of bending rifidity and tilt angle")))
annotate_figure(plots,
  top=text_grob(title))
```

**Figure 1:**



Both the bending and the tilt variable seem to have a lot of outliers when looking at figure 1. Even when

transformed using log10 they still seem to be quite skewed. This is something we need to keep in mind in future research.

## Exploring relations between variables

### Density plot

I will now make density plots of some of the variables to see whether or not they look promising in finding class distinctions. These density plots will also make any skewing in the data apparent.

```
# Create density plots
p <- ggplot(nona_data, aes(x=bending, colour = factor(sterol.type))) +
  geom_density() +
  xlab(get_des_by_ab(codebook, "bending")) +
  ylab("Density") +
  ggtitle("Bending rigidity density plot") +
  labs(colour = get_des_by_ab(codebook, "sterol.type"))

p2 <- ggplot(nona_data, aes(x=APL, colour = factor(sterol.type))) +
  geom_density() +
  xlab(get_des_by_ab(codebook, "APL")) +
  ylab("Density") +
  ggtitle("Area per lipid density plot") +
  labs(colour = get_des_by_ab(codebook, "sterol.type"))

p3 <- ggplot(nona_data, aes(x=compress, colour = factor(sterol.type))) +
  geom_density() +
  xlab(get_des_by_ab(codebook, "compress")) +
  ylab("Density") +
  ggtitle("Compressibility density plot") +
  labs(colour = get_des_by_ab(codebook, "sterol.type"))

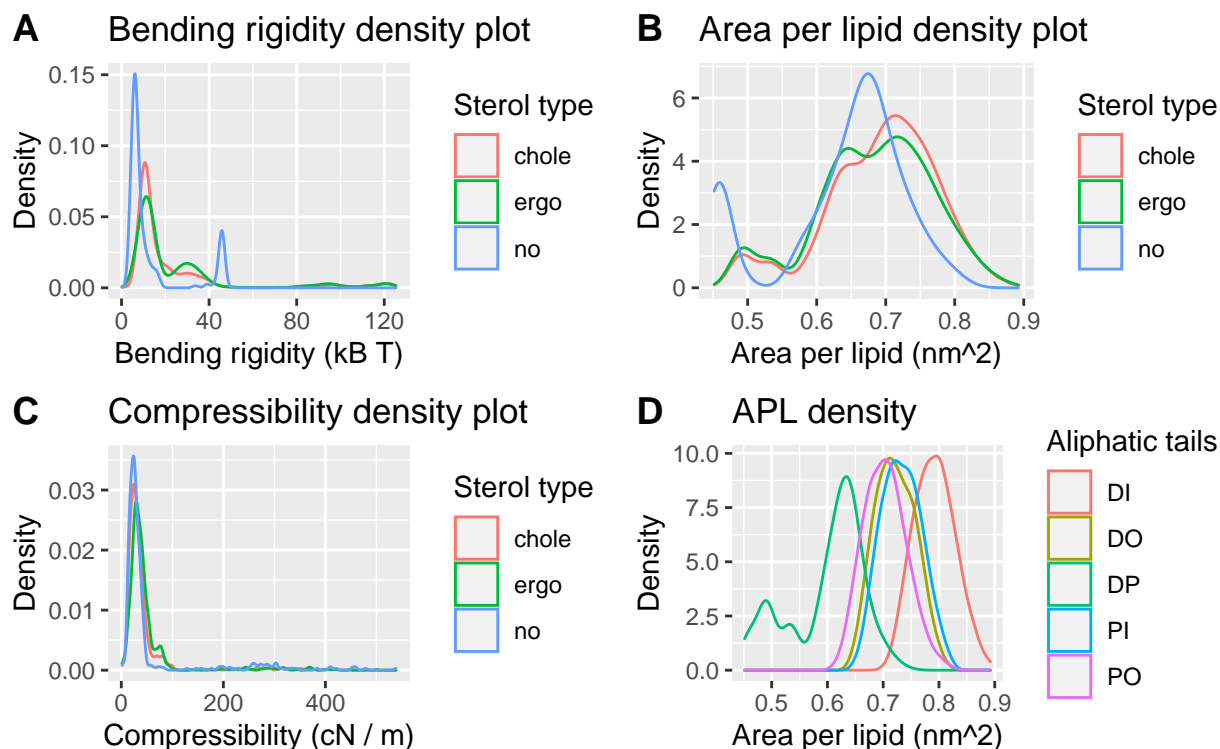
p4 <- ggplot(nona_data, aes(x=APL, colour = factor(tails))) +
  geom_density() +
  xlab(get_des_by_ab(codebook, "APL")) +
  ylab("Density") +
  ggtitle("APL density") +
  labs(colour = get_des_by_ab(codebook, "tails"))

# Plot density plots
plots <- ggarrange(p, p2, p3, p4,
  labels = c("A", "B", "C", "D"),
  ncol = 2, nrow = 2)
title <- expression(atop(bold("Figure 2:"),
  scriptstyle(paste("Density plots of multiple",
    " variables coloured on sterol type"))))
annotate_figure(plots,
  top=text_grob(title))
```



**Figure 2:**

Density plots of multiple variables coloured on sterol type



Plot A in figure 2 show that there is basically no distinction between cholesterol and ergosterol based on the bending rigidity. However, distinction between no sterol or a sterol does seem to be possible based on the bending rigidity seeing as the peaks of there classes only overlaps a small bit. There does seem to be an odd peak in the no sterol class at around 50 kB/t bending rigidity.

In plot B it looks like there is absolutely no way of distinguishing different classes based on the area per lipid, seeing as the peaks are basically in the same place. But just to be certain we'll still plot in against some other variables to make sure, seeing as it might be a very useful variable when paired with something like bending rigidity. The results depicted in plot C also don't look very promising seeing as every peak is around the same place again. Plot D on the other hand seems a lot more interesting seeing as all of the peaks are in slightly different locations, the DO, PO and the PI seem to overlap a lot but the DP and DI tails seem to have little overlap with the rest. It seems like APL could possibly be used as a variable to distinguish between aliphatic tails when paired with another variable.

When looking at all of the plots it looks like most of the data is skewed in one way or another, even though when we looked at the structure there didn't seem to be much. This can probably be explained by the extra dimension that was added in these plots, the sterol type and the aliphatic tails.

## Scatter plot

```
# Code not run
# Results are too big and too clumpy for the analysis
pairs(nona_data,
      panel = "panel.smooth",
      pch = 20,
```

```
cex = 0.35,
col = rgb(0.1, 0.4, 1, alpha = 0.4))
```

This code is for a paired scatter plot. The eval is currently set to false seeing as there are a lot of variables that are compared to each other in this data set. Running this code will result in a big plot that isn't very well interpretable in a pdf file.

First we'll plot the bending area per lipid against the rigidity of the membrane, colouring the points based on the sterol type, to look for any kind of clustering or correlation. To make it even clearer we'll also make a second plot with a loess regression. The reason we use loess regression instead of linear regression is because loess will show us a better trend line, seeing as it isn't trying to fit a straight line like linear regression.

```
#Plot APL against bending rigidity
chart <- ggplot(nona_data, aes(APL, bending, colour = factor(sterol.type) )) +
  geom_point(alpha=1/10, size=0.5) +
  xlab(get_des_by_ab(codebook, "APL")) +
  ylab(get_des_by_ab(codebook, "bending")) +
  theme(axis.title.y = element_text(size=8)) +
  ggtitle("APL against bending rigidity scatter plot") +
  scale_color_manual(name=get_des_by_ab(codebook, "sterol.type"),
                    labels = c("Cholesterol",
                              "Ergosterol",
                              "No sterol"),
                    values = hue_pal()(3))

plots <- ggarrange(chart,
  chart +
    geom_smooth(method = "loess") + ggtitle(""),
  labels = c("A", "B"),
  ncol = 1, nrow = 2)
```

```
## 'geom_smooth()' using formula 'y ~ x'
```


```
title <- expression(atop(bold("Figure 3:"),
  scriptstyle(paste("APL against bending rigidity,",
                    " with and without loess regression")))))

# Give the plot a number
annotate_figure(plots,
  top=text_grob(title))
```


**Figure 3:**

APL against bending rigidity, with and without loess regression

**A** APL against bending rigidity scatter plot



**B**



Looking at the plots in figure 3 we can see that there is a clear separation in the area per lipid to bending rigidity ratio when comparing no sterol present to both cholesterol and ergosterol. When comparing cholesterol and ergosterol we can't see such a clear separation, there seems to be a lot of overlap. There are also a few outliers that are barely visible at area per lipid 0.4 and bending rigidity 0. The correlation seems to be inverse logarithmic, so to check this I'll make a plot where the bending rigidity is log transformed.

```
# Create plot
chart <- ggplot(nona_data, aes(APL, log10(bending + 1), colour = factor(sterol.type) )) +
  geom_point(alpha=1/10, size=0.5) +
  xlab(get_des_by_ab(codebook, "APL")) +
  ylab(paste( "Log10 ", get_des_by_ab(codebook, "bending"))) +
  geom_smooth(method = "loess") +
  ggtitle("APL against log10 bending rigidity scatter plot") +
  scale_color_manual(name=get_des_by_ab(codebook, "sterol.type"),
    labels = c("Cholesterol",
               "Ergosterol",
               "No sterol"),
    values = hue_pal()(3))

# Plot the Plot and add a tag
chart + labs(tag = "Figure 4: APL against log transformed bending rigidity") +
  theme(plot.title = element_text(hjust = 0.5),
    plot.margin = margin(t = 10, r = 10, b = 40, l = 10),
    plot.tag.position = c(0.4, -0.1)
  )

## 'geom_smooth()' using formula 'y ~ x'
```




Figure 4: APL against log transformed bending rigidity

Looking at figure 4 we can see that there indeed seem to be an almost linear correlation between the log10 transformed bending rigidity and the area per lipid. There does still seem to be a slight curve in all of the trend lines when the area per lipid increases.

Here I'll plot the temperature against the bending rigidity while colouring based on saturation index, and deciding the shape based on the sterol type.

```
# Create rectangles to fit over the plot
rects <- data.frame(xstart = seq(0,600,300), xend = seq(300,900,300), col = letters[1:3])

# Create the plot
chart <- ggplot(nona_data, aes(temperature, bending, colour = factor(satur.index),
                              shape = factor(sterol.type) )) +
  geom_jitter(alpha=1/10) +
  xlab(get_des_by_ab(codebook, "temperature")) +
  ylab(get_des_by_ab(codebook, "bending")) +
  ggtitle("Temperature against bending rigidity,
          coloured by saturation index, shaped on sterol type") +
  theme(plot.title = element_text(size=10)) +
  labs(colour = get_des_by_ab(codebook, "sterol.type"),
       shape = get_des_by_ab(codebook, "satur.index"))

# Print the plot and add 2 boxes for the 2 groups in the temperature variable
chart +
  geom_rect(aes(xmin = -Inf, xmax = 312.9, ymin = -Inf, ymax = Inf),
           color="orange",
           fill = "green", alpha = 0.0001) +
```

```
geom_rect(aes(xmin = 313.1, xmax = Inf, ymin = -Inf, ymax = Inf),
          color="purple",
          fill = "yellow", alpha = 0.0001) +
labs(tag = "Figure 5: Temperature against bending rigidity") +
theme(plot.title = element_text(hjust = 0.5),
      plot.margin = margin(t = 10, r = 10, b = 40, l = 10),
      plot.tag.position = c(0.35, -0.1)
    )
```




Figure 5: Temperature against bending rigidity

When looking at the plot in figure 5, there seem to be a clear separation between a saturation index of 0 while the rest of the values seem to be clustered quite close together. There does seem to be a small separation between 0.5, 1 and 2 values. An interesting observation is that there don't seem to be any data point with a saturation index of more than 0.5 in the 328 kelvin group (the purple box).

Here the APL will be plot against the bending rigidity again, but here the colour is decided based on the saturation index and the shape on the sterol type.

```
# Create plots
chart2 <- ggplot(nona_data, aes(APL, bending, colour = factor(satur.index), shape = factor(sterol.type))
  geom_jitter(alpha=1, size=0.5) +
  xlab(get_des_by_ab(codebook, "APL")) +
  ylab(get_des_by_ab(codebook, "bending")) +
  ggtitle("APL against bending rigidity,
           coloured on saturation index, shaped on sterol type") +
  theme(plot.title = element_text(size=10))
chart2 + labs(colour = get_des_by_ab(codebook, "satur.index"), shape = get_des_by_ab(codebook, "sterol.type"),
  labs(tag = "Figure 6: APL against bending rigidity coloured on saturation index") +
```

```
theme(plot.title = element_text(hjust = 0.5),
      plot.margin = margin(t = 10, r = 10, b = 40, l = 10),
      plot.tag.position = c(0.4, -0.1)
    )
```

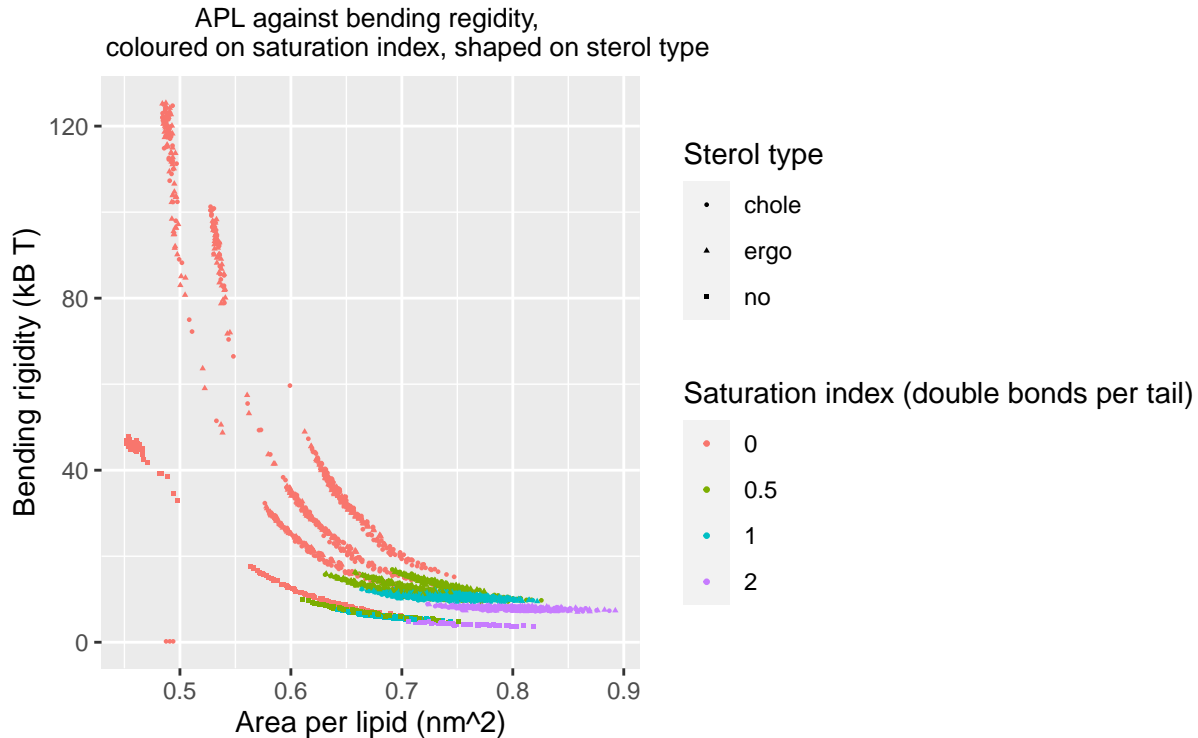


Figure 6: APL against bending rigidity coloured on saturation index

In figure 6 we can see the same kind of correlation as before, but now we can see that there also seem to be some kind of clustering based on the saturation index.

Here is a plot that shows the compressibility plotted against the bending rigidity.

```
# Create plot
chart <- ggplot(nona_data, aes(compress, bending, colour = factor(sterol.type) )) +
  geom_point(alpha=2/10, size=0.5) +
  xlab(get_des_by_ab(codebook, "compress")) +
  ylab(get_des_by_ab(codebook, "bending")) +
  theme(axis.title.y = element_text(size=8)) +
  ggtitle("Compressibility against bending rigidity,
          coloured on sterol type") +
  scale_color_manual(name=get_des_by_ab(codebook, "sterol.type"),
                    labels = c("Cholesterol",
                              "Ergosterol",
                              "No sterol"),
                    values = hue_pal()(3)) +
  theme(plot.title = element_text(size=10))

chart2 <- ggplot(nona_data, aes(compress, bending, colour = factor(sterol.type) )) +
  geom_point(alpha=2/10, size=0.5) +
```

```

xlab(get_des_by_ab(codebook, "compress")) +
ylab(get_des_by_ab(codebook, "bending")) +
theme(axis.title.y = element_text(size=8)) +
ggtitle("Compressibility against bending rigidity,
        coloured on sterol type, zoomed in") +
scale_color_manual(name=get_des_by_ab(codebook, "sterol.type"),
                   labels = c("Cholesterol",
                              "Ergosterol",
                              "No sterol"),
                   values = hue_pal()(3)) +
scale_x_continuous(limits=c(0,120)) +
scale_y_continuous(limits=c(0, 45)) +
theme(plot.title = element_text(size=10))

# Arrange plots
plots <- ggarrange(chart,
                   chart2,
                   labels = c("A", "B"),
                   ncol = 1, nrow = 2)

```

## Warning: Removed 261 rows containing missing values (geom\_point).

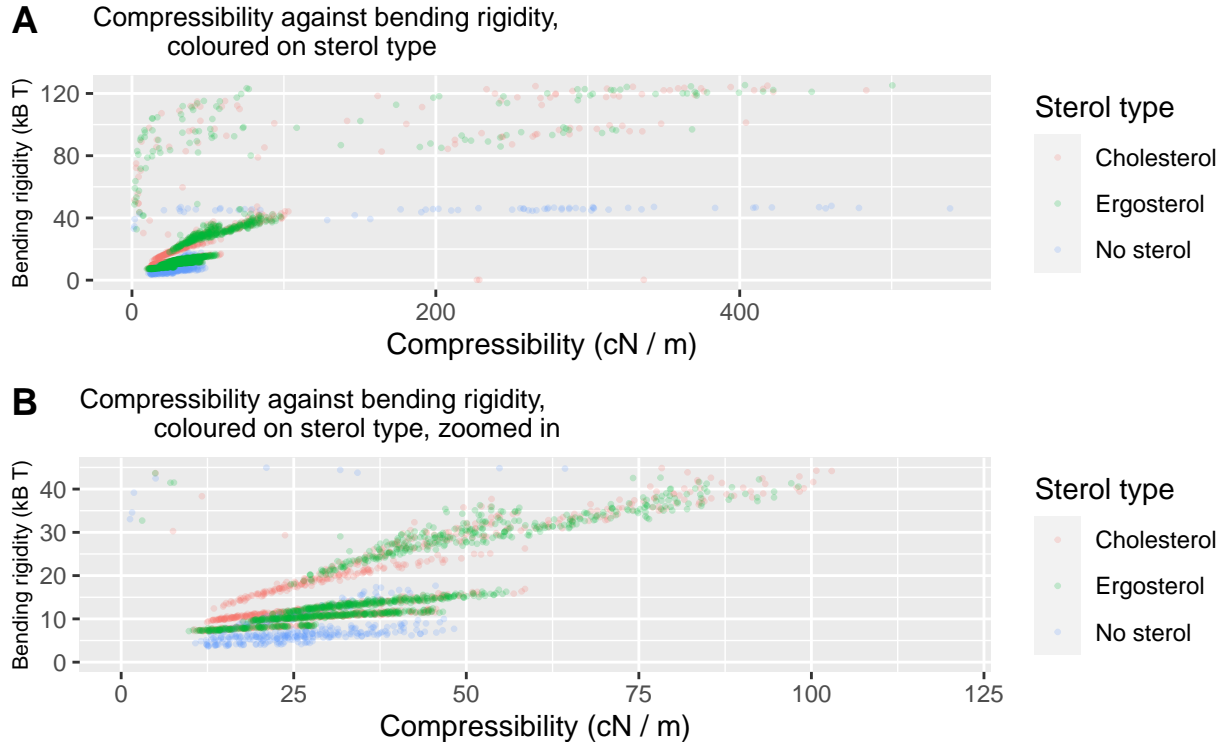
```

title <- expression(atop(bold("Figure 7:"),
                          scriptstyle("Compressibility against bending rigidity, full and zoomed"))))
annotate_figure(plots,
                top=text_grob(title))

```

**Figure 7:**

Compressibility against bending rigidity, full and zoomed



In plot A of figure 7 we can't really make out much of a pattern in the dense clout to the left, but when looking at the rest of the data points there does seem to be some grouping based on whether or not no sterol or a sterol is present. When zooming in in plot B we can now see the same kind of grouping happen again based on either a sterol present or no sterol present.

## Heatmap

Here a heatmap is made to look for correlation between variables.

```
# Get correlation
correlation <- nona_data[,9:14]
melted_cormat <- melt(cor(correlation))

ggplot(data = melted_cormat, aes(x=Var1, y=Var2, fill=value)) +
  geom_tile() +
  ggtitle("Heatmap of the variables") +
  labs(tag = "Figure 8: Heatmap of the independent variables") +
  theme(plot.title = element_text(hjust = 0.5),
        plot.margin = margin(t = 10, r = 10, b = 40, l = 10),
        plot.tag.position = c(0.5, -0.1)
  )
```






Figure 8: Heatmap of the independent variables

The results shown in figure 8 are very interesting. It seems like all of the variables seem to be quite heavily correlated, except for APL, which is negatively correlated. It also seem like the comprehensibility is slightly less correlated then the rest of the variables. But thickness, bending, tilt and z order all seem to be heavily correlated.

## PCA

Here we'll make a PCA plot to look for possible clusters.

```
# PCA plot results
pca_res <- prcomp(correlation, scale. = TRUE, center = TRUE)
print(pca_res)

## Standard deviations (1, ..., p=6):
## [1] 2.24944886 0.69482659 0.55396674 0.31774787 0.20132585 0.09391948
##
## Rotation (n x k) = (6 x 6):
##          PC1      PC2      PC3      PC4      PC5
## APL      -0.4039636 -0.38324528 -0.4127355 -0.705249607 0.14566451
## thickness 0.4238589 0.31048244 0.2156260 -0.511891894 -0.01429709
## bending  0.4152014 -0.14463396 -0.5644389 0.026515944 -0.69807879
## tilt     0.4196997 -0.08623976 -0.5085343 0.252592319 0.68632375
## zorder   0.4334754 0.22183399 0.0793969 -0.419528906 0.14134816
## compress 0.3474421 -0.82410989 0.4468250 -0.008899588 0.01561576
##          PC6
```

```
## APL          0.031475796
## thickness    0.644369420
## bending      -0.008970044
## tilt         0.151540515
## zorder       -0.748732960
## compress     0.012300032
```

```
# Create PCA plot
```

```
plots <- ggarrange(autoplot(pca_res, data = nona_data,
                           colour = 'sterol.type', alpha=2/10,
                           size = 5/10) +
  ggtitle("PCA plot") +
  scale_color_manual(name=get_des_by_ab(codebook, "sterol.type"),
                    labels = c("Cholesterol",
                              "Ergosterol",
                              "No sterol"),
                    values = hue_pal()(3)),
  autoplot(pca_res, data = nona_data,
           colour = 'sterol.type', alpha=1/5,
           size = 1/2) +
  ggtitle("PCA plot zoomed in") +
  scale_color_manual(name=get_des_by_ab(codebook, "sterol.type"),
                    labels = c("Cholesterol",
                              "Ergosterol",
                              "No sterol"),
                    values = hue_pal()(3)) +
  scale_x_continuous(limits=c(-0.03,0.025)) +
  scale_y_continuous(limits=c(-0.03,0.04)),
  labels = c("A", "B"),
  ncol = 1, nrow = 2)
```

```
## Warning: Removed 271 rows containing missing values (geom_point).
```


```
# Add title
```

```
title <- expression(atop(bold("Figure 9:"),
  scriptstyle("PCA plot using autoplot"))))
annotate_figure(plots,
  top=text_grob(title))
```


**Figure 9:**

PCA plot using autoplot

**A** PCA plot



**B** PCA plot zoomed in



In plot A of figure 9 we can hardly see any clusters at all, it does seem to have a dense cloud to the right. Plot B cuts off the data point on the side and zooms in on the dense cloud of dots visible in plot A, a total of 271 dots aren't visible because of this. In this plot we can't really make out much of a pattern or cluster. It does, however, seem like no sterol cluster is slightly separated from the other sterol types.

Here is another way to make a PCA plot that is a bit more advanced seeing as it will also give us a way to identify correlation based on variables in the clustering.

```
# Create PCA plot
res.pca <- prcomp(correlation, scale = TRUE)
plot1 <- fviz_eig(res.pca) +
  theme(axis.title.y = element_text(size=6))
plot2 <- fviz_pca_var(res.pca,
  col.var = "contrib", # Colour by contributions to the PC
  gradient.cols = c("orange", "purple"),
  repel = TRUE        # Avoid text overlapping
)
plot3 <- fviz_pca_biplot(res.pca, repel = FALSE,
  col.var = "purple", # Variables colour
  col.ind = nona_data$sterol.type, # Individuals colour
  label = FALSE,
  addEllipses = TRUE,
  ellipse.type = "t"
)

# Arrange plots
plot <- arrangeGrob(plot1, plot2,
  plot3,
```

```

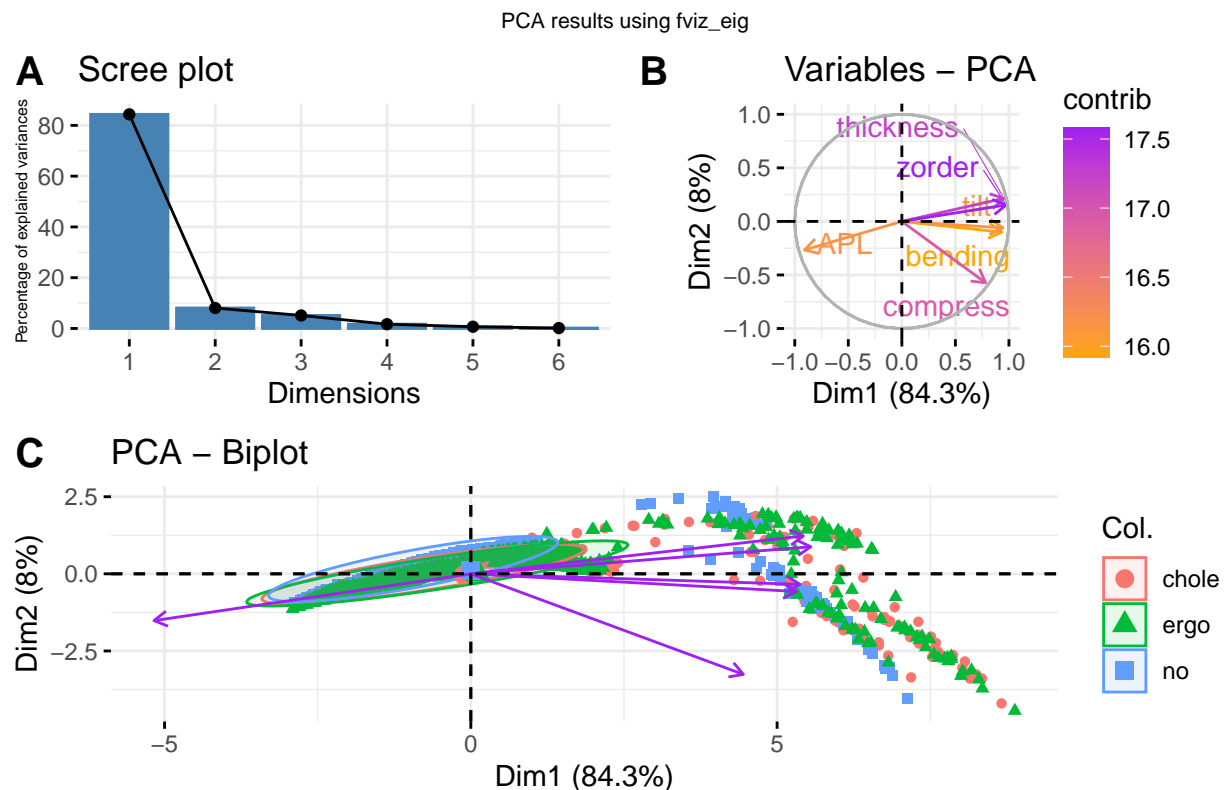
ncol = 2, nrow = 2,
layout_matrix = rbind(c(1,2), c(3,3)))

## Warning in MASS::cov.trob(data[, vars]): Probable convergence failure

plot <- as_ggplot(plot) +
  draw_plot_label(label = c("A", "B", "C"), size = 15,
    x = c(0, 0.5, 0), y = c(1, 1, 0.5))
title <- expression(atop(bold("Figure 10:"),
  scriptstyle("PCA results using fviz_eig")))
annotate_figure(plot,
  top=text_grob(title))

```

**Figure 10:**



The PCA plots, B and C, shown in figure 10 confirms what we could already see in the heat map, that is that most of the variables seem to be heavily correlated, seeing as they are pointing in the same direction, except for APL and compress. We can also see that thickness and z order are correlated just like tilt and bending. The PCA plot also still looks the same, only now it has as ellipse for clusters that was calculated using a t test. The ellipses all seem to overlap so there doesn't seem to be much clustering.

## K means

Here we'll use K-means clustering as a different technique to find possible clusters.

```

# Calculate K-means, set seed for reproducibility
set.seed(666)
km.res <- kmeans(scale(correlation), 3, nstart = 25)

# Dimension reduction using PCA
res.pca <- prcomp(correlation, scale = TRUE)
ind.coord <- as.data.frame(get_pca_ind(res.pca)$coord)
ind.coord$cluster <- factor(km.res$cluster)
ind.coord$sterol.type <- nona_data$sterol.type
kable(head(ind.coord), caption="kmeans results") %>%
  kable_styling(latex_options = c("scale_down", "hold_position"))

```

Table 4: kmeans results

Dim.1	Dim.2	Dim.3	Dim.4	Dim.5	Dim.6	cluster	sterol.type
-1.385930	-0.1910350	-0.1129858	0.0445117	0.0120431	0.0276672	3	chole
-1.468677	-0.2591053	-0.1694398	0.0073676	0.0181114	-0.0044196	3	chole
-1.331954	-0.1460367	-0.0861045	0.0548778	0.0031295	0.0500533	3	chole
-1.709264	-0.4213788	-0.3268945	-0.0624653	0.0414548	-0.1137487	3	chole
-1.205414	-0.0421895	-0.0158859	0.0909438	-0.0115402	0.1168805	3	chole
-1.610915	-0.3573446	-0.2555837	-0.0356559	0.0387487	-0.0660767	3	chole

The results in table 4 are printed to give a quick overview of the results of the K-means and to check for any possible problems. The row with sterol type is added so we can look whether or not this has anything to do with the found clusters.

```

# Percentage of variance explained by dimensions
eigenvalue <- round(get_eigenvalue(res.pca), 1)
variance.percent <- eigenvalue$variance.percent
kable(head(eigenvalue), caption="Variance explained by dimensions") %>%
  kable_styling(latex_options = c("scale_down", "hold_position"))

```

Table 5: Variance explained by dimensions

	eigenvalue	variance.percent	cumulative.variance.percent
Dim.1	5.1	84.3	84.3
Dim.2	0.5	8.0	92.4
Dim.3	0.3	5.1	97.5
Dim.4	0.1	1.7	99.2
Dim.5	0.0	0.7	99.9
Dim.6	0.0	0.1	100.0

In table 5 we can see the amount of variance is explained by ever dimension. Just like in the previous PCA plot we can see that over 80% of the variance is explained by a single variable, which is odd, to say the least.

```
# Make pca plot based on K-means clusters
ggscatter(
  ind.coord, x = "Dim.1", y = "Dim.2",
  color = "cluster", palette = "npg", ellipse = TRUE, ellipse.type = "convex",
  shape = "sterol.type", size = 1.5, legend = "right", ggtheme = theme_bw(),
  xlab = paste0("Dim 1 (", variance.percent[1], "% )"),
  ylab = paste0("Dim 2 (", variance.percent[2], "% )")
) +
  stat_mean(aes(color = cluster), size = 4) +
  labs(tag = "Figure 11: PCA plot based on K-means clustering",
       shape = get_des_by_ab(codebook, "sterol.type")) +
  theme(plot.title = element_text(hjust = 0.5),
        plot.margin = margin(t = 10, r = 10, b = 40, l = 10),
        plot.tag.position = c(0.46, -0.1)
  )
```




Figure 11: PCA plot based on K-means clustering

This plot shown in figure 11 still doesn't seem very promising seeing as there seem to be three clear clusters, but this seems to be more because of the large spread of the data. The clusters also seem to have nothing to do with the sterol type of the membrane.

## Research question

Is it possible to use machine learning to reliably predicts the sterol type in a membrane with a higher than 80% accuracy, given the area per lipid, the bending rigidity and the compressibility?

## Data Cleaning


Here is a quick code snippet to show us the skewness of the data. This will be useful when deciding whether or not to use standardization or normalization on the columns.

```
apply(nona_data[c("temperature","sterol.conc","satur.index",
                  "PC.conc","ethanol.conc","APL", "thickness", "bending",
                  "tilt", "zorder","compress")], 2, skewness, na.rm =TRUE)
```

```
##   temperature   sterol.conc   satur.index   PC.conc   ethanol.conc
## 1.1257170406 -0.1719109112 0.9075759350 0.0008402866 0.0149496771
##          APL   thickness   bending   tilt   zorder
## -0.6069614189 0.7652531496 2.6743851360 2.6321330960 1.1463852300
##      compress
## 3.9327898578
```

A lot of the data seems to be extremely skewed. Lets plot it in histograms to get a better look.

```
hist.data.frame(nona_data[c("temperature","sterol.conc","satur.index",
                             "PC.conc","ethanol.conc","APL", "thickness", "bending",
                             "tilt", "zorder","compress")])
```




The data seems to be skewed a lot in almost every column. If we were planning on using a machine learning algorithms that are based on the assumption that the data has a Gaussian distribution we would need to use normalization (transformation) to force the data in to such a distribution. To achieve this we could use a log

or cube root transformation. We probably won't be using any algorithms that assume that our distribution is Gaussian, which is why we won't make any transformations at first. Distributions like bending, tilt and compress have very long tails due to outliers. These could be log transformed. However they do not follow a power law distribution (they aren't fat tailed, the tail is just long), and values do not span across multiple orders of magnitude. This is why they won't be log transformed. Seeing as most of the distributions aren't Gaussian we should use min\_max normalization (scaling). This does lessen/suppress the impact of outliers, it also isn't as effective with outliers. We could also use z score normalization (standardization), but this assumes that our distributions are Gaussian, which they are not. It could still be used. However, due to the heavy skew and the high amount of outliers in some column, the scale probably won't be balanced. We could also look at box cox transformation, but I wasn't taught this method, so I'm not confident about using it. The reason we use normalization is because we might want to use an algorithm like k nearest neighbor (KNN), which is an algorithm that is very sensitive to data scale.

```
scale_min_max <- function(x) (x - min(x)) / (max(x) - min(x))
nona_copy <- nona_data[c("APL", "thickness", "bending",
                        "tilt", "zorder", "compress")]

scaled_data <- as.data.frame(lapply(nona_copy, FUN=scale_min_max))
hist.data.frame(scaled_data)
```



```
# scaled_data <- as.data.frame(lapply(nona_copy, FUN=scale))
# hist.data.frame(scaled_data)

apply(scaled_data, 2, skewness, na.rm = TRUE)
```

```
##          APL  thickness  bending      tilt  zorder  compress
```



```
## -0.6069614  0.7652531  2.6743851  2.6321331  1.1463852  3.9327899
```

After min max scaling we can see that the distributions stay the same but scale on all of the variables now ranges from zero to one. In the future we could always try different forms of normalization and transformation to see whether or not our model improves. Now the only thing left to do is add the sterol type label.

```
scaled_data$sterol.type <- nona_data$sterol.type
kable(head(scaled_data), caption="Cleaned data") %>%
  kable_styling(latex_options = c("hold_position"), full_width = TRUE)
```

APL	thickness	bending	tilt	zorder	compress	sterol.type
0.6474103	0.2983812	0.0789085	0.0361181	0.1640910	0.0426240	chole
0.6685354	0.2831408	0.0793886	0.0340809	0.1603082	0.0428898	chole
0.6361482	0.3096582	0.0803025	0.0371130	0.1674086	0.0418494	chole
0.7216122	0.2350388	0.0784531	0.0290836	0.1506323	0.0403502	chole
0.6079452	0.3370028	0.0819857	0.0406205	0.1715037	0.0406188	chole
0.6999791	0.2554618	0.0771624	0.0313995	0.1544628	0.0420252	chole

```
# Write new data to csv
write.csv(scaled_data, "data/clean_data.csv", row.names = FALSE)
```

Only the sterol type was added as the label, seeing as the other variables were considered labels by the assignment, and we're interested in the sterol type in the membrane. Looking at the PCA plot and the heatmap it would seem like some of the variables are heavily correlated. However, it was decided not to remove any of the variables as of this time.

Here is a PCA plot that is made using the data after applying log and cube transformations to reduce skewness.

```
# Create PCA plot
new_help <- log10(correlation[2:6])
new_help$APL <- correlation$APL^2
apply(new_help, 2, skewness, na.rm = TRUE)
```

```
## thickness bending tilt zorder compress APL
## 0.5275196 0.7523384 1.2834053 -0.4950526 1.0032745 -0.2422874
```

```
res.pca <- prcomp(new_help, scale = TRUE)
plot1 <- fviz_eig(res.pca) +
  theme(axis.title.y = element_text(size=6))
plot2 <- fviz_pca_var(res.pca,
  col.var = "contrib", # Colour by contributions to the PC
  gradient.cols = c("orange", "purple"),
  repel = TRUE # Avoid text overlapping
)
plot3 <- fviz_pca_biplot(res.pca, repel = FALSE,
  col.var = "purple", # Variables colour
  col.ind = nona_data$sterol.type, # Individuals colour
  label = FALSE,
  addEllipses = TRUE,
```

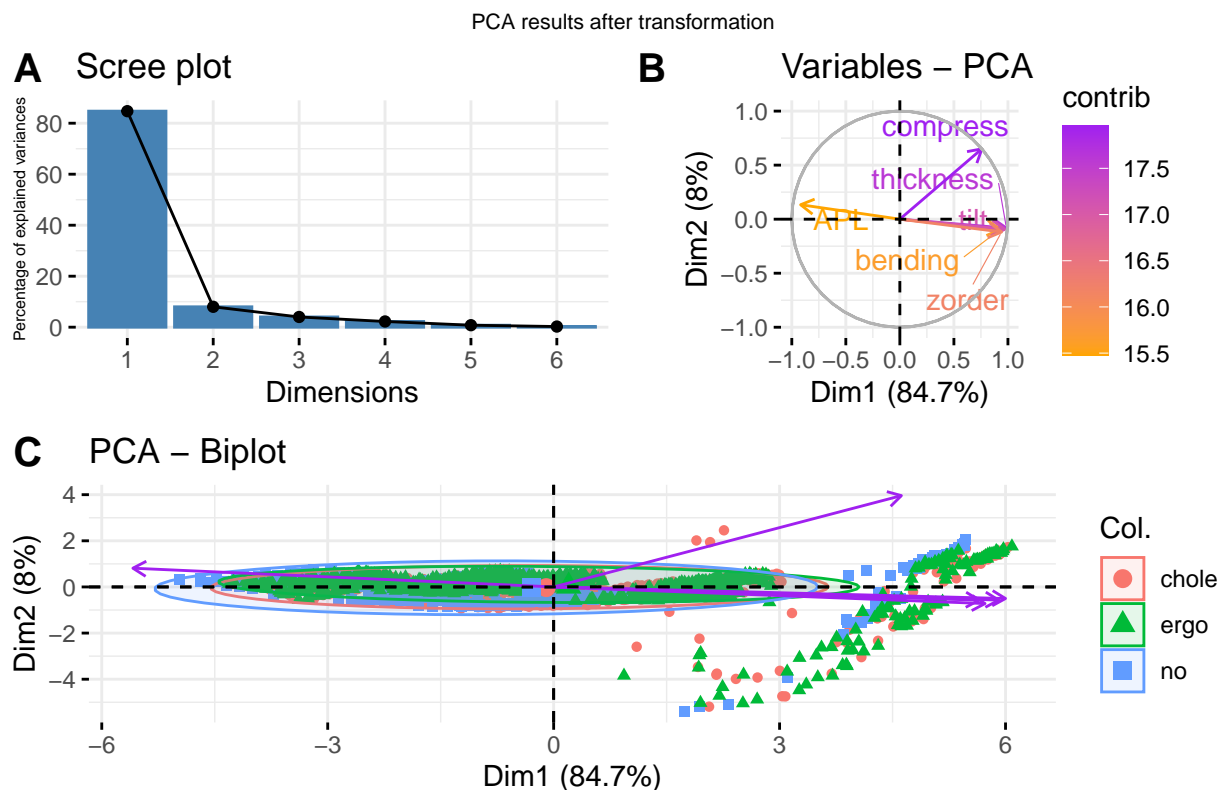
```

        ellipse.type = "t"
      )
# Arrange plots
plot <- arrangeGrob(plot1, plot2,
                    plot3,
                    ncol = 2, nrow = 2,
                    layout_matrix = rbind(c(1,2), c(3,3)))

plot <- as_ggplot(plot) +
  draw_plot_label(label = c("A", "B", "C"), size = 15,
                 x = c(0, 0.5, 0), y = c(1, 1, 0.5))
title <- expression(atop(bold("Figure 14:"),
                          scriptstyle("PCA results after transformation"))
                    )
annotate_figure(plot,
               top=text_grob(title))

```

**Figure 14:**



As we can see in figure 14, reducing the skewness had absolutely no effect, besides increasing the variance explained by the first dimension and flipping the PCA upside down.

```

kable(head(data), caption="Original data") %>%
  kable_styling(latex_options = c("scale_down", "hold_position"))

```

```

kable(head(scaled_data), caption="Cleaned data") %>%
  kable_styling(latex_options = c("scale_down", "hold_position"))

```

Table 7: Original data

temperature	sterol.type	sterol.conc	other.phosph	tails	satur.index	PC.conc	ethanol.conc	APL	thickness	bending	tilt	zorder	compress
298	chole	20	PE	PI	1	33	20	0.736911	3.628973	10.08020	14.7578	0.175499	24.18218
298	chole	20	PE	PI	1	50	20	0.746226	3.596991	10.14030	14.3075	0.172084	24.32495
298	chole	20	PE	PI	1	25	20	0.731945	3.652637	10.25470	14.9777	0.178494	23.76607
298	chole	20	PE	PI	1	100	20	0.769630	3.496052	10.02320	13.2029	0.163349	22.96077
298	chole	20	PE	PI	1	0	20	0.719509	3.710018	10.46540	15.7530	0.182191	23.10503
298	chole	20	PE	PI	1	75	20	0.760091	3.538908	9.86163	13.7148	0.166807	23.86050

Table 8: Cleaned data

APL	thickness	bending	tilt	zorder	compress	sterol.type
0.6474103	0.2983812	0.0789085	0.0361181	0.1640910	0.0426240	chole
0.6685354	0.2831408	0.0793886	0.0340809	0.1603082	0.0428898	chole
0.6361482	0.3096582	0.0803025	0.0371130	0.1674086	0.0418494	chole
0.7216122	0.2350388	0.0784531	0.0290836	0.1506323	0.0403502	chole
0.6079452	0.3370028	0.0819857	0.0406205	0.1715037	0.0406188	chole
0.6999791	0.2554618	0.0771624	0.0313995	0.1544628	0.0420252	chole

## Weka

### Rescaling data

Here we rescale all of the possible labels to factors so when it is exported to a .arff file it will be encoded correctly.

```
scale_min_max <- function(x) (x - min(x)) / (max(x) - min(x))
nona_copy <- nona_data
```

```
#str(nona_copy)
```

```
nona_copy[1:8] <- lapply(nona_copy[1:8], factor)
str(nona_copy)
```

```
## tibble [2,690 x 14] (S3: tbl_df/tbl/data.frame)
## $ temperature : Factor w/ 2 levels "298","328": 1 1 1 1 1 1 1 1 1 ...
## $ sterol.type : Factor w/ 3 levels "chole","ergo",...: 1 1 1 1 1 1 1 1 1 ...
## $ sterol.conc : Factor w/ 4 levels "0","10","20",...: 3 3 3 3 3 3 3 3 3 ...
## $ other.phosph: Factor w/ 1 level "PE": 1 1 1 1 1 1 1 1 1 ...
## $ tails       : Factor w/ 5 levels "DI","D0","DP",...: 4 4 4 4 4 4 2 2 2 ...
## $ satur.index : Factor w/ 4 levels "0","0.5","1",...: 3 3 3 3 3 3 3 3 3 ...
## $ PC.conc     : Factor w/ 7 levels "0","25","33",...: 3 4 2 7 1 6 5 7 4 3 ...
## $ ethanol.conc: Factor w/ 7 levels "0","5","10","15",...: 5 5 5 5 5 5 5 2 2 2 ...
## $ APL         : num [1:2690] 0.737 0.746 0.732 0.77 0.72 ...
## $ thickness   : num [1:2690] 3.63 3.6 3.65 3.5 3.71 ...
## $ bending     : num [1:2690] 10.1 10.1 10.3 10 10.5 ...
## $ tilt        : num [1:2690] 14.8 14.3 15 13.2 15.8 ...
## $ zorder      : num [1:2690] 0.175 0.172 0.178 0.163 0.182 ...
## $ compress    : num [1:2690] 24.2 24.3 23.8 23 23.1 ...
## - attr(*, "na.action")= 'omit' Named int [1:153] 281 400 401 402 403 404 405 406 414 415 ...
## ..- attr(*, "names")= chr [1:153] "281" "400" "401" "402" ...
```

```

index <- sapply(nona_copy, is.numeric)
nona_copy[index] <- lapply(nona_copy[index], scale_min_max)

#scaled_data <- as.data.frame(lapply(nona_copy, FUN=scale_min_max))
#hist.data.frame(scaled_data)
#str(nona_copy)

```

Here is a function to put a specific column as the last column for the encoding in the .arff file. Some of the code is commented out so it won't create a lot of datasets that take up memory.

```

library(RWeka)
Mutate_df <- function(data, last_wanted_column){
  data %>%
    relocate(all_of(last_wanted_column), .after = last_col()) %>%
    select(tail(names(.), 7)) %>%
    #mutate_at(last_wanted_column, as.factor) %>%
    write.arff(file = paste0("data/scaled_data_", last_wanted_column, ".arff"))
}
Mutate_df(nona_copy, "tails")
#Mutate_df(nona_copy, "sterol.type")
#Mutate_df(nona_copy, "sterol.conc")

```

Here is some code to make a .arff file.

```

library(RWeka)
write.arff(scaled_data, file = "data/scaled_data.arff")

```

## Quality metrics

The two most important metrics that we'll be using are the accuracy and the confusion matrix. The accuracy is important because it tells us how accurate the algorithm is. The confusion matrix is also very important seeing as this can tell us if the model may be prediction some classes with high accuracy while have a very poor accuracy on other classes.

## Output

### Sterol prediction

```

Tester: weka.experiment.PairedCorrectedTTester -G 4,5,6 -D 1 -R 2 -S 0.05 -result-matrix
"weka.experiment.ResultMatrixLatex -mean-prec 2 -stddev-prec 0 -col-name-width 0 -row-name-width 0
-mean-width 2 -stddev-width 2 -sig-width 1 -count-width 0 -print-row-names -enum-col-names"
Analysing: Percent_correct
Datasets: 1
Resultsets: 9
Confidence: 0.05 (two tailed)
Sorted by: -
Date: 10/10/21, 4:18 PM

```

These aren't the result we were hoping for. It seems like there is no way of discerning between sterol type, seeing as the highest achieved accuracy is equal to around 61%. This is quite poor to say the least. Maybe we should take a closer look at the confusion matrix of some of the algorithms to see if we can find out what's going on here.

Table 9: Sterol prediction accuracy

Dataset	(1)	(2)	(3)	(4)	(5)	(6)	(7)	(8)	(9)
R-data-frame	42.94	56.65 ◦	59.51 ◦	47.83 ◦	57.30 ◦	60.20 ◦	57.07 ◦	58.03 ◦	70.03 ◦

◦, • statistically significant improvement or degradation

Table 10: Sterol prediction accuracy (Key)

- (1) rules.ZeroR " 48055541465867954
- (2) rules.OneR '-B 50' -3459427003147861443
- (3) lazy.IBk '-K 100 -W 0 -A \\'weka.core.neighboursearch.LinearNNSearch -A \\'weka.core.EuclideanDistance -R first-last\\\' -3080186098777067172
- (4) bayes.NaiveBayes " 5995231201785697655
- (5) trees.J48 '-C 0.25 -M 200' -21773316839364444
- (6) trees.RandomForest '-P 100 -I 100 -num-slots 1 -K 0 -M 1.0 -V 0.001 -S 1' 1116839470751428698
- (7) functions.SMO '-C 1.0 -L 0.001 -P 1.0E-12 -N 0 -V -1 -W 1 -K \\'functions.supportVector.PolyKernel -E 1.0 -C 250007\' -calibrator \\'functions.Logistic -R 1.0E-8 -M -1 -num-decimal-places 4\' -6585883636378691736
- (8) functions.SimpleLogistic '-I 0 -M 500 -H 50 -W 0.0' 7397710626304705059
- (9) trees.LMT '-I -1 -M 15 -W 0.0' -1113212459618104943

## Tails prediction

Tester: weka.experiment.PairedCorrectedTTTester -G 4,5,6 -D 1 -R 2 -S 0.05 -result-matrix

"weka.experiment.ResultMatrixLatex -mean-prec 2 -stddev-prec 0 -col-name-width 0 -row-name-width 0 -mean-width 2 -stddev-width 2 -sig-width 1 -count-width 0 -print-row-names -enum-col-names"

Analysing: Percent\_correct

Datasets: 1

Resultsets: 9

Confidence: 0.05 (two tailed)

Sorted by: -

Date: 10/10/21, 4:26 PM

Table 11: Accuracy tails prediction

Dataset	(1)	(2)	(3)	(4)	(5)	(6)	(7)	(8)	(9)
R-data-frame	37.81	82.28 ◦	82.39 ◦	77.83 ◦	81.03 ◦	99.47 ◦	79.71 ◦	99.42 ◦	99.94 ◦

◦, • statistically significant improvement or degradation

## Temporary results

After running a multitude of test in weka to try and make a model to predict sterol type, I found something quite interesting. It seems that it is almost impossible to predict the sterol type in a membrane. By this I mean it is very hard to distinguish between ergosterol and cholesterol. It is however completely possible to predict whether or not there is a sterol present in the membrane. After running some test with the all of the instances with no sterol removed we see that almost every algorithm go's to 50% accuracy, meaning it's basically as good as zeroR. Even when using voting with OneR, ZeroR, RandomForest, J48, NaiveBayes, SimpleLogistic and SMO, we still only get 55.7006 % correctly classified instances (with the no sterol group removed). This gives us an interesting conclusion to the temporary research question. No, it is not possible to predict the sterol in a membrane given certain measurements. It is however possible to distinguish between no sterol and a sterol. This leads me to conclude that there is just very little to no difference between ergosterol and cholesterol, seeing as when testing to predict the sterol concentration we get a 99.1297% using the voting of the same algorithms as used on sterol type (again on the data with instances with no sterol removed), and only given the measurements.

Table 12: Accuracy tails prediction (Key)

- (1) rules.ZeroR " 48055541465867954
- (2) rules.OneR '-B 50' -3459427003147861443
- (3) lazy.IBk '-K 100 -W 0 -A \"weka.core.neighboursearch.LinearNNSearch -A \\\\\"weka.core.EuclideanDistance -R first-last\\\\\"' -3080186098777067172
- (4) bayes.NaiveBayes " 5995231201785697655
- (5) trees.J48 '-C 0.25 -M 200' -217733168393644444
- (6) trees.RandomForest '-P 100 -I 100 -num-slots 1 -K 0 -M 1.0 -V 0.001 -S 1' 1116839470751428698
- (7) functions.SMO '-C 1.0 -L 0.001 -P 1.0E-12 -N 0 -V -1 -W 1 -K \"functions.supportVector.PolyKernel -E 1.0 -C 250007\" -calibrator \"functions.Logistic -R 1.0E-8 -M -1 -num-decimal-places 4\"' -6585883636378691736
- (8) functions.SimpleLogistic '-I 0 -M 500 -H 50 -W 0.0' 7397710626304705059
- (9) trees.LMT '-I -1 -M 15 -W 0.0' -1113212459618104943

## EDA Discussion

I need to make this very clear. The EDA was made with a specific question in mind. This probably exposed the exploration steps to an subconscious bias when making plots and exploring relations between variables.